

Scrapping the Web

kgold for 11/3

HTML

- "HyperText Markup Language": **markup** as opposed to **programming** (instructions for rendering text, but not general programs)
- Instructions that aren't supposed to be rendered appear in **tags** - often these surround text to indicate which text is affected
 - `Bold`
 - `<i>Italics</i>`
 - `Anchor tags for links`

Structure of a classic minimal web page

```
<html>
<head>
<title> My amazing webpage!</title>
</head>
<body>
<p>Some text for the page; p stands for paragraph.</p>

<p>Next paragraph, let's link <a href="http://www.google.com">somewhere.</a>

</body>
</html>
```

Moving ahead in time: CSS

- Rather than attempt to do everything in HTML, web designers delineate logical containers for content, and give rules for rendering it elsewhere
 - The containers are often `<div>` tags with named classes; the style rules affect the named classes
 - The style rules are kept in a separate `.css` file that says how to render everything
- This makes it difficult for an automated parser to reason about where text is throughout the page, without actually rendering the page

Sample page snippet using CSS

`<body>`

`<!-- Primary Page Layout`

```
<div class="container">
  <div class="row">
    <div class="main-container column">
      <div class="logo-wrapper">
        <a href="index.html"></a>
      </div>
      <nav>
        <ul>
          <li><a href="about.html">About</a></li>
          <li><a href="publications.html">Publications</a></li>
          <li><a href="games.html">Games</a></li>
        </ul>
      </nav>
    </div>
  </div>
</div>
```

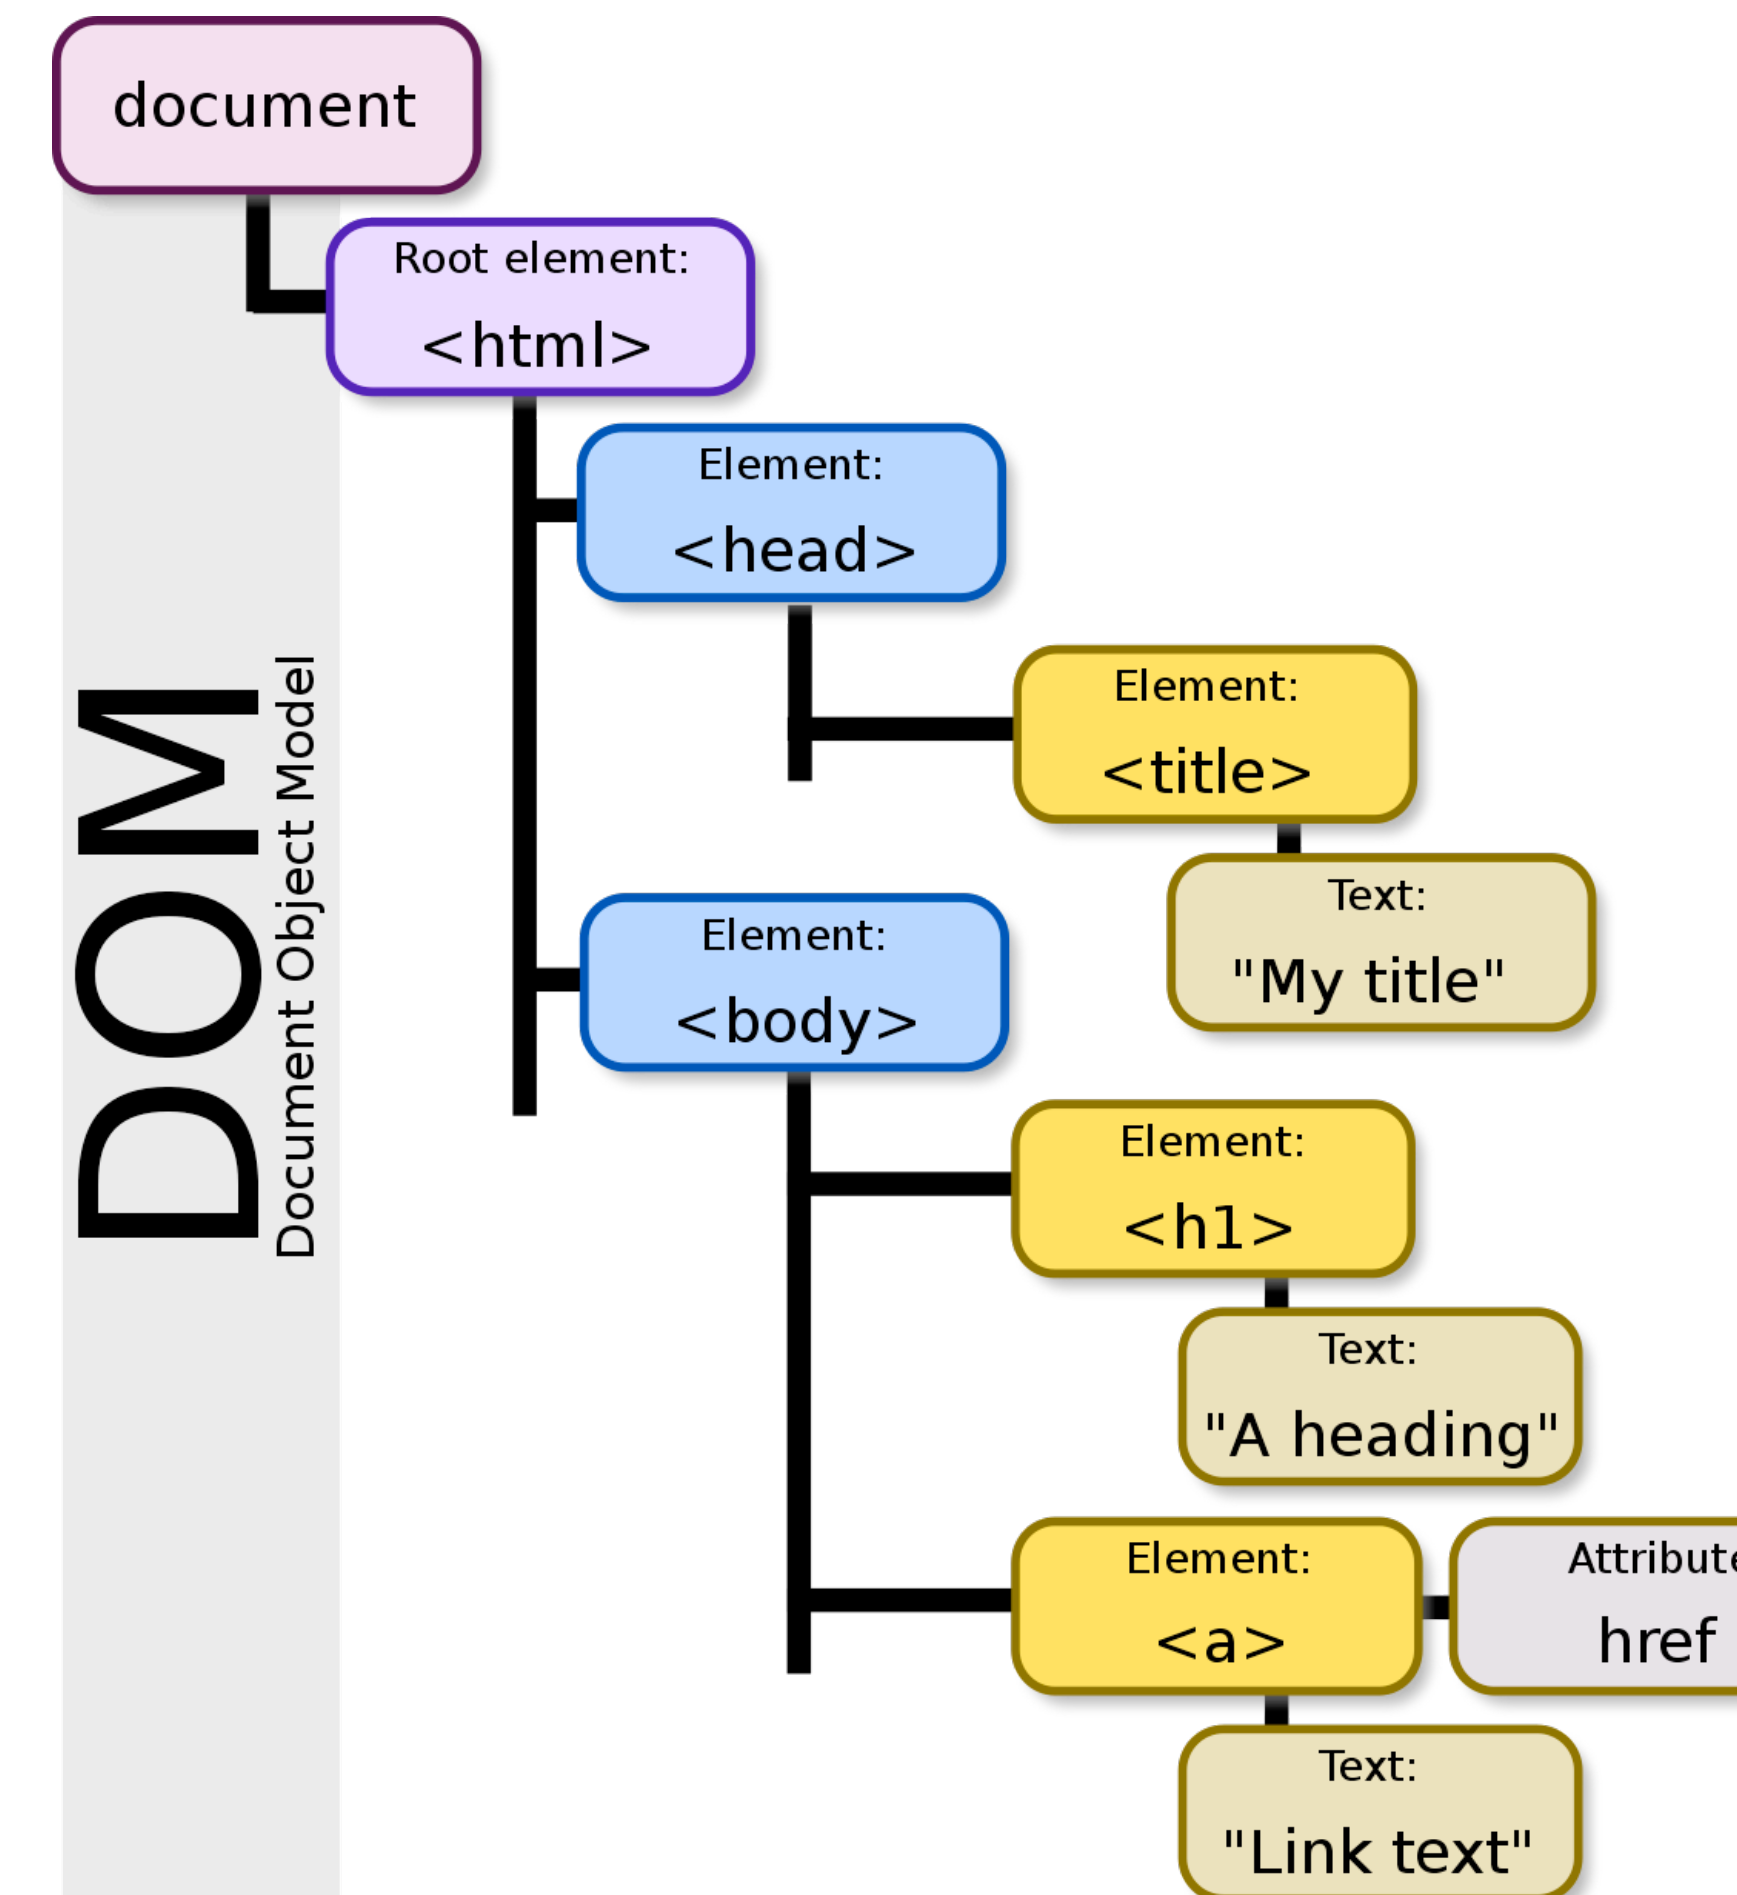
HTML fragment

```
nav {
  background: rgba(131, 120, 108, 0.7);
}
nav ul {
  list-style: none;
  text-align: center;
  padding: 0;
  margin: 0;
}
nav li {
  display: inline-block;
  margin: 0;
}
nav a {
  color: #fff;
  text-decoration: none;
  padding: 1rem;
  display: block;
}
```

CSS fragment

JavaScript and the DOM

- The more interactive webpages became, the more they used stronger programming languages to be interactive - JavaScript is most common
- These languages can require execution to build out the webpage - the page could **no longer be readable by a scraping bot**
- JavaScript interacts with a tree-structured model of the page called the DOM: Document Object Model
 - Parts of the tree could be dynamically generated or deleted in response to user actions



Getting a Webpage

- Unlike proprietary APIs, webpages are served to everyone without needing permission first
 - Although some have policies against "robots" visiting them, and possibly even safeguards against them
- Pages are served on port 80, where the server listens for GET requests and hands over the appropriate content
- The HTML and any accompanying CSS and JavaScript are served as text for the browser to interpret
 - Images, audio, and video could also be served as binary files

Getting a Webpage with requests

- requests: An easy module for getting the text of webpages

Included in Colab

```
import requests
page = requests.get('http://www.bu.edu')
contents = page.content
```

contents is HTML as string

What you get in 2021 (www.bu.edu)

[illegible]

Other features of the requests module

- Access sites needing authentication (if you have access)

```
>>> r = requests.get('https://api.github.com/user', auth=('user', 'pass'))
```

- Pass in arguments to website (like search query) as a dict

```
>>> payload = {'key1': 'value1', 'key2': ['value2', 'value3']}  
  
>>> r = requests.get('https://httpbin.org/get', params=payload)  
>>> print(r.url)  
https://httpbin.org/get?key1=value1&key2=value2&key2=value3
```

- Retrieve binary data, not just webpages

```
from PIL import Image  
from io import BytesIO  
  
i = Image.open(BytesIO(r.content))
```

Beautiful Soup 4

- This module is for parsing HTML, retrieving what text it can find and reconstructing the tree
- Two basic use cases are:
 - Retrieving all outgoing links
 - Retrieving all the text meant for human consumption

Creating the BeautifulSoup (tree) with a parser

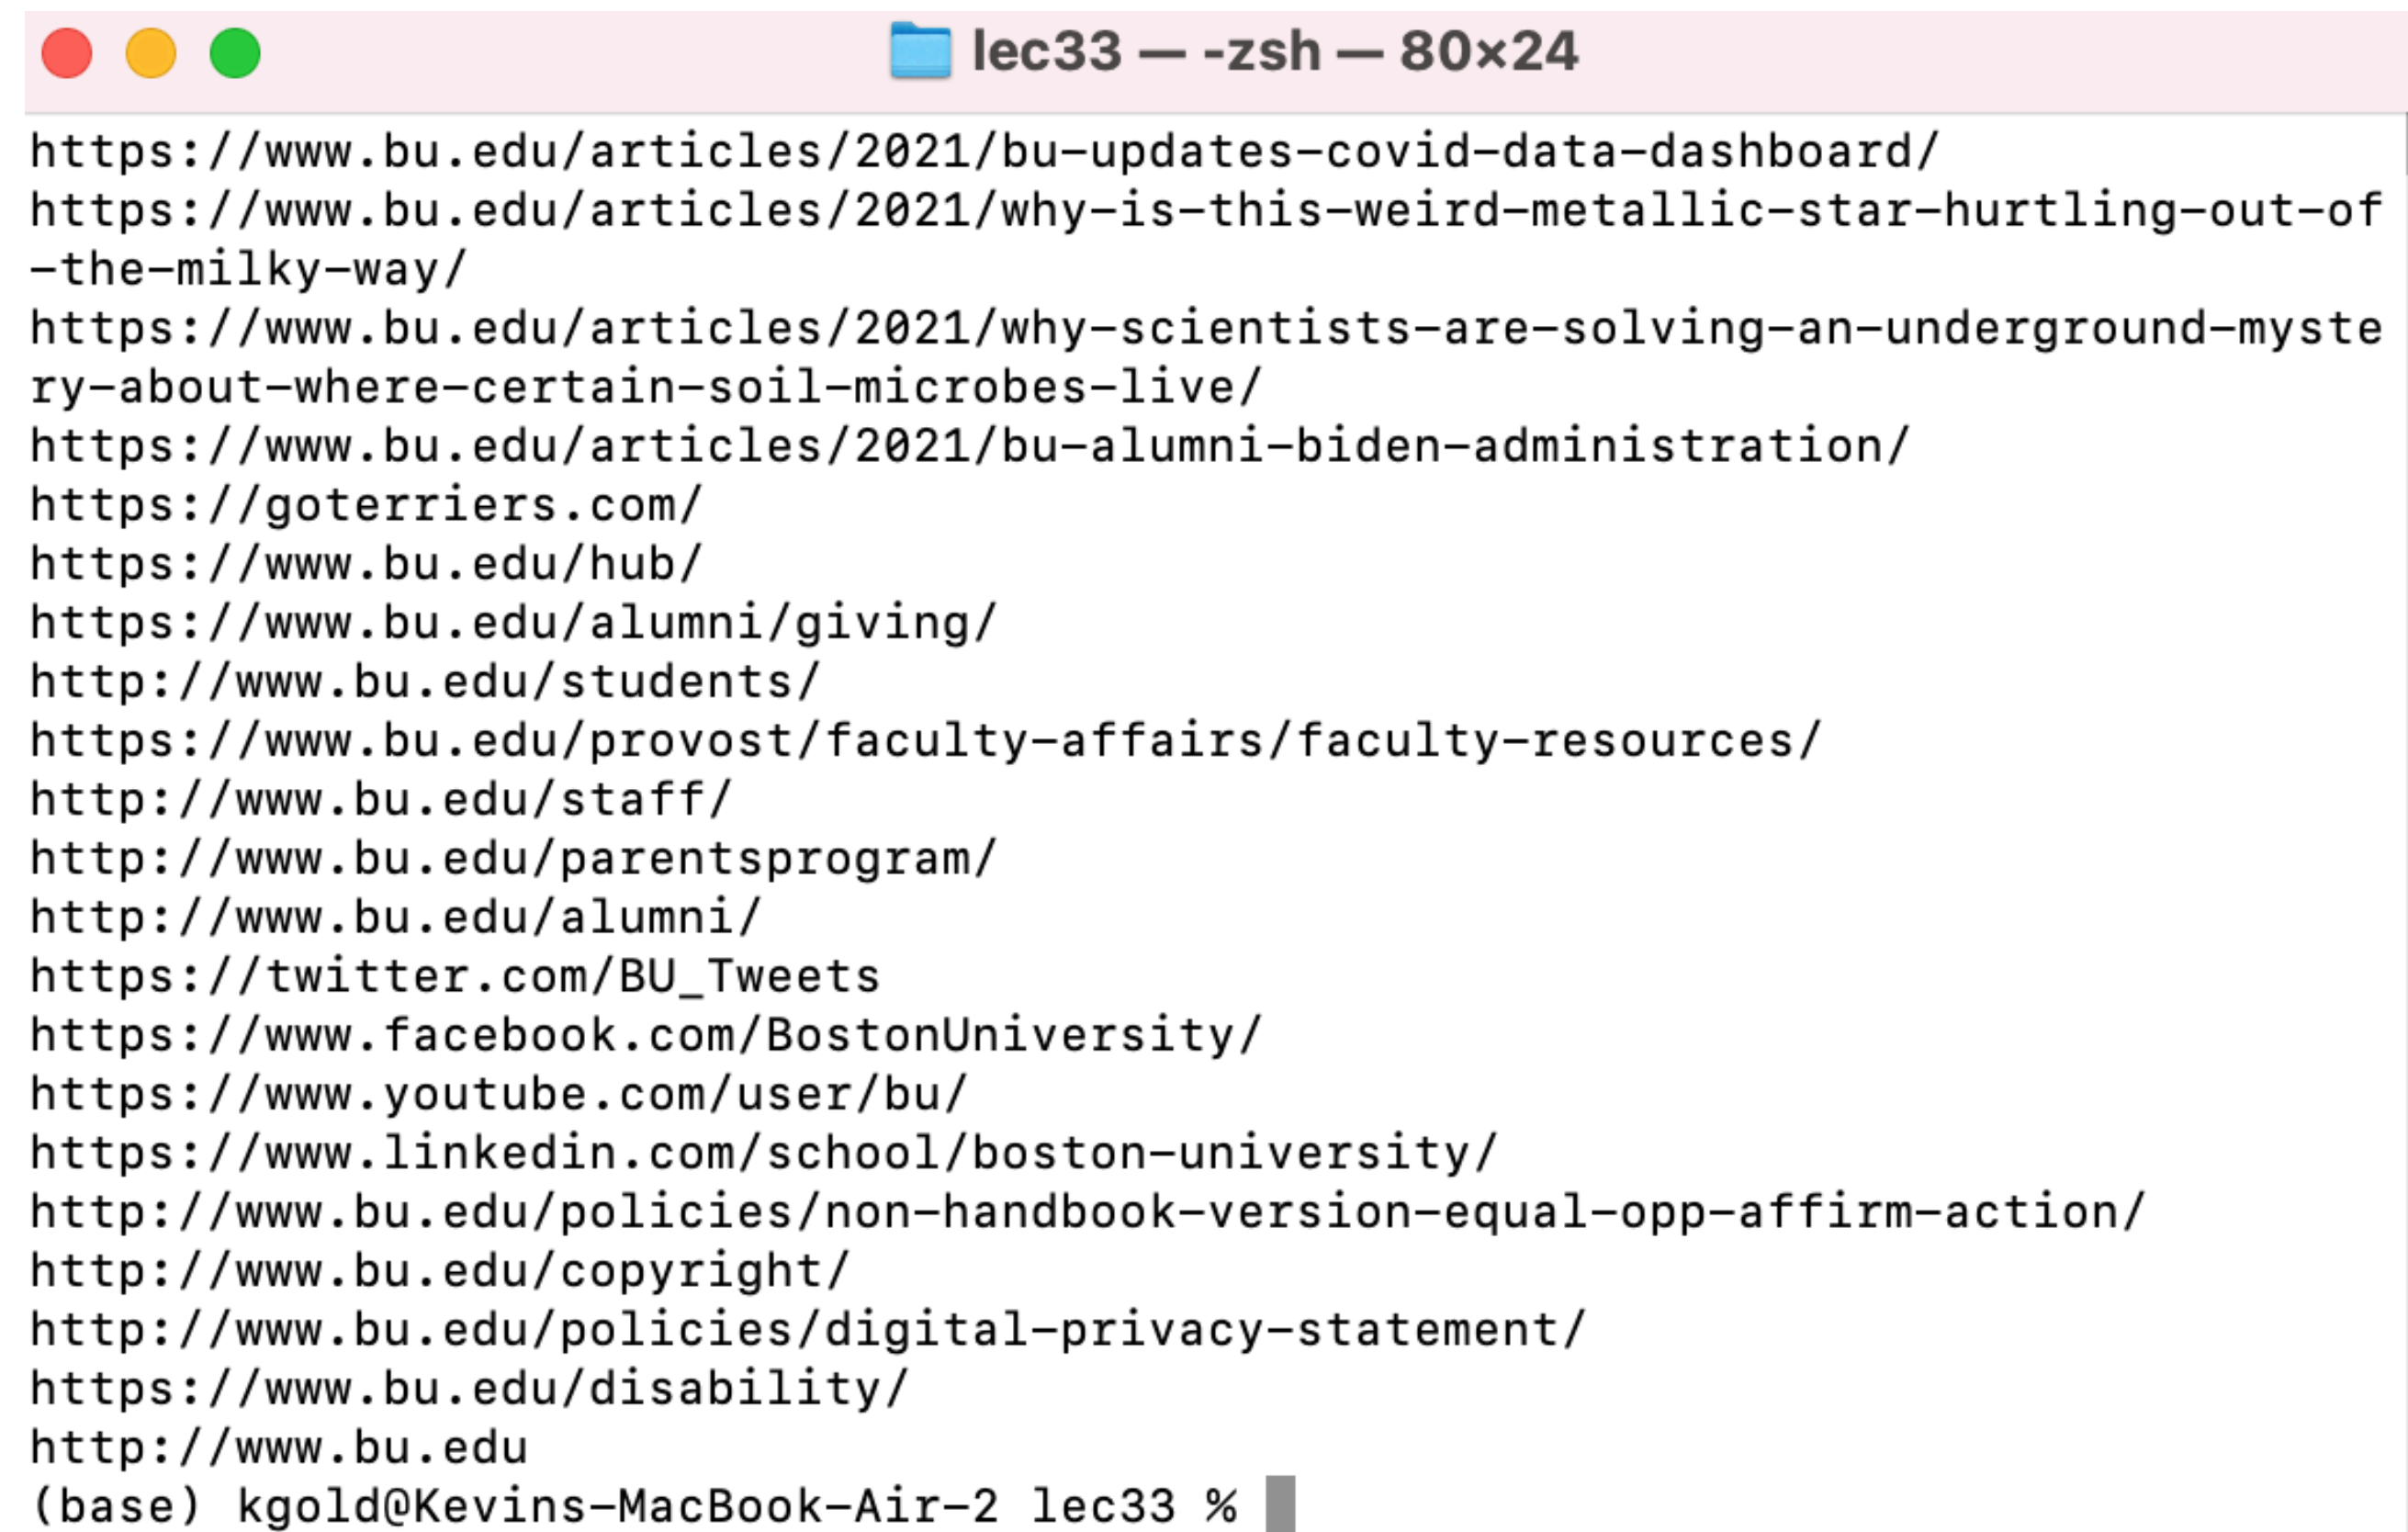
Included in colab

```
from bs4 import BeautifulSoup
soup = BeautifulSoup(html_doc, 'html.parser')

print(soup.prettify())
# <html>
# <head>
# <title>
#   The Dormouse's story
# </title>
# </head>
# <body>
# <p class="title">
#   <b>
#     The Dormouse's story
#   </b>
# </p>
# <p class="story">
#   Once upon a time there were three little sisters; and their names were
#   <a class="sister" href="http://example.com/elsie" id="link1">
#     Elsie
#   </a>
#
```

Finding all the links

```
for link in soup.find_all('a'):  
    print(link.get('href'))
```




A terminal window titled "lec33 — -zsh — 80x24" displays the output of a BeautifulSoup script. The output is a list of 25 URLs found on the Boston University website. The URLs include various university pages, social media links, and external sites like goterriers.com. The terminal shows the output line by line, with a cursor at the bottom.

```
https://www.bu.edu/articles/2021/bu-updates-covid-data-dashboard/  
https://www.bu.edu/articles/2021/why-is-this-weird-metallic-star-hurtling-out-of  
-the-milky-way/  
https://www.bu.edu/articles/2021/why-scientists-are-solving-an-underground-myste  
ry-about-where-certain-soil-microbes-live/  
https://www.bu.edu/articles/2021/bu-alumni-biden-administration/  
https://goterriers.com/  
https://www.bu.edu/hub/  
https://www.bu.edu/alumni/giving/  
http://www.bu.edu/students/  
https://www.bu.edu/provost/faculty-affairs/faculty-resources/  
http://www.bu.edu/staff/  
http://www.bu.edu/parentsprogram/  
http://www.bu.edu/alumni/  
https://twitter.com/BU_Tweets  
https://www.facebook.com/BostonUniversity/  
https://www.youtube.com/user/bu/  
https://www.linkedin.com/school/boston-university/  
http://www.bu.edu/policies/non-handbook-version-equal-opp-affirm-action/  
http://www.bu.edu/copyright/  
http://www.bu.edu/policies/digital-privacy-statement/  
https://www.bu.edu/disability/  
http://www.bu.edu  
(base) kgold@Kevins-MacBook-Air-2 lec33 %
```

Finding all the links

```
for link in soup.find_all('a'):  
    print(link.get('href'))
```


- Sample applications:
 - Find highly interconnected networks of pages that form "communities" to understand a page's ecosystem



Finding all the links

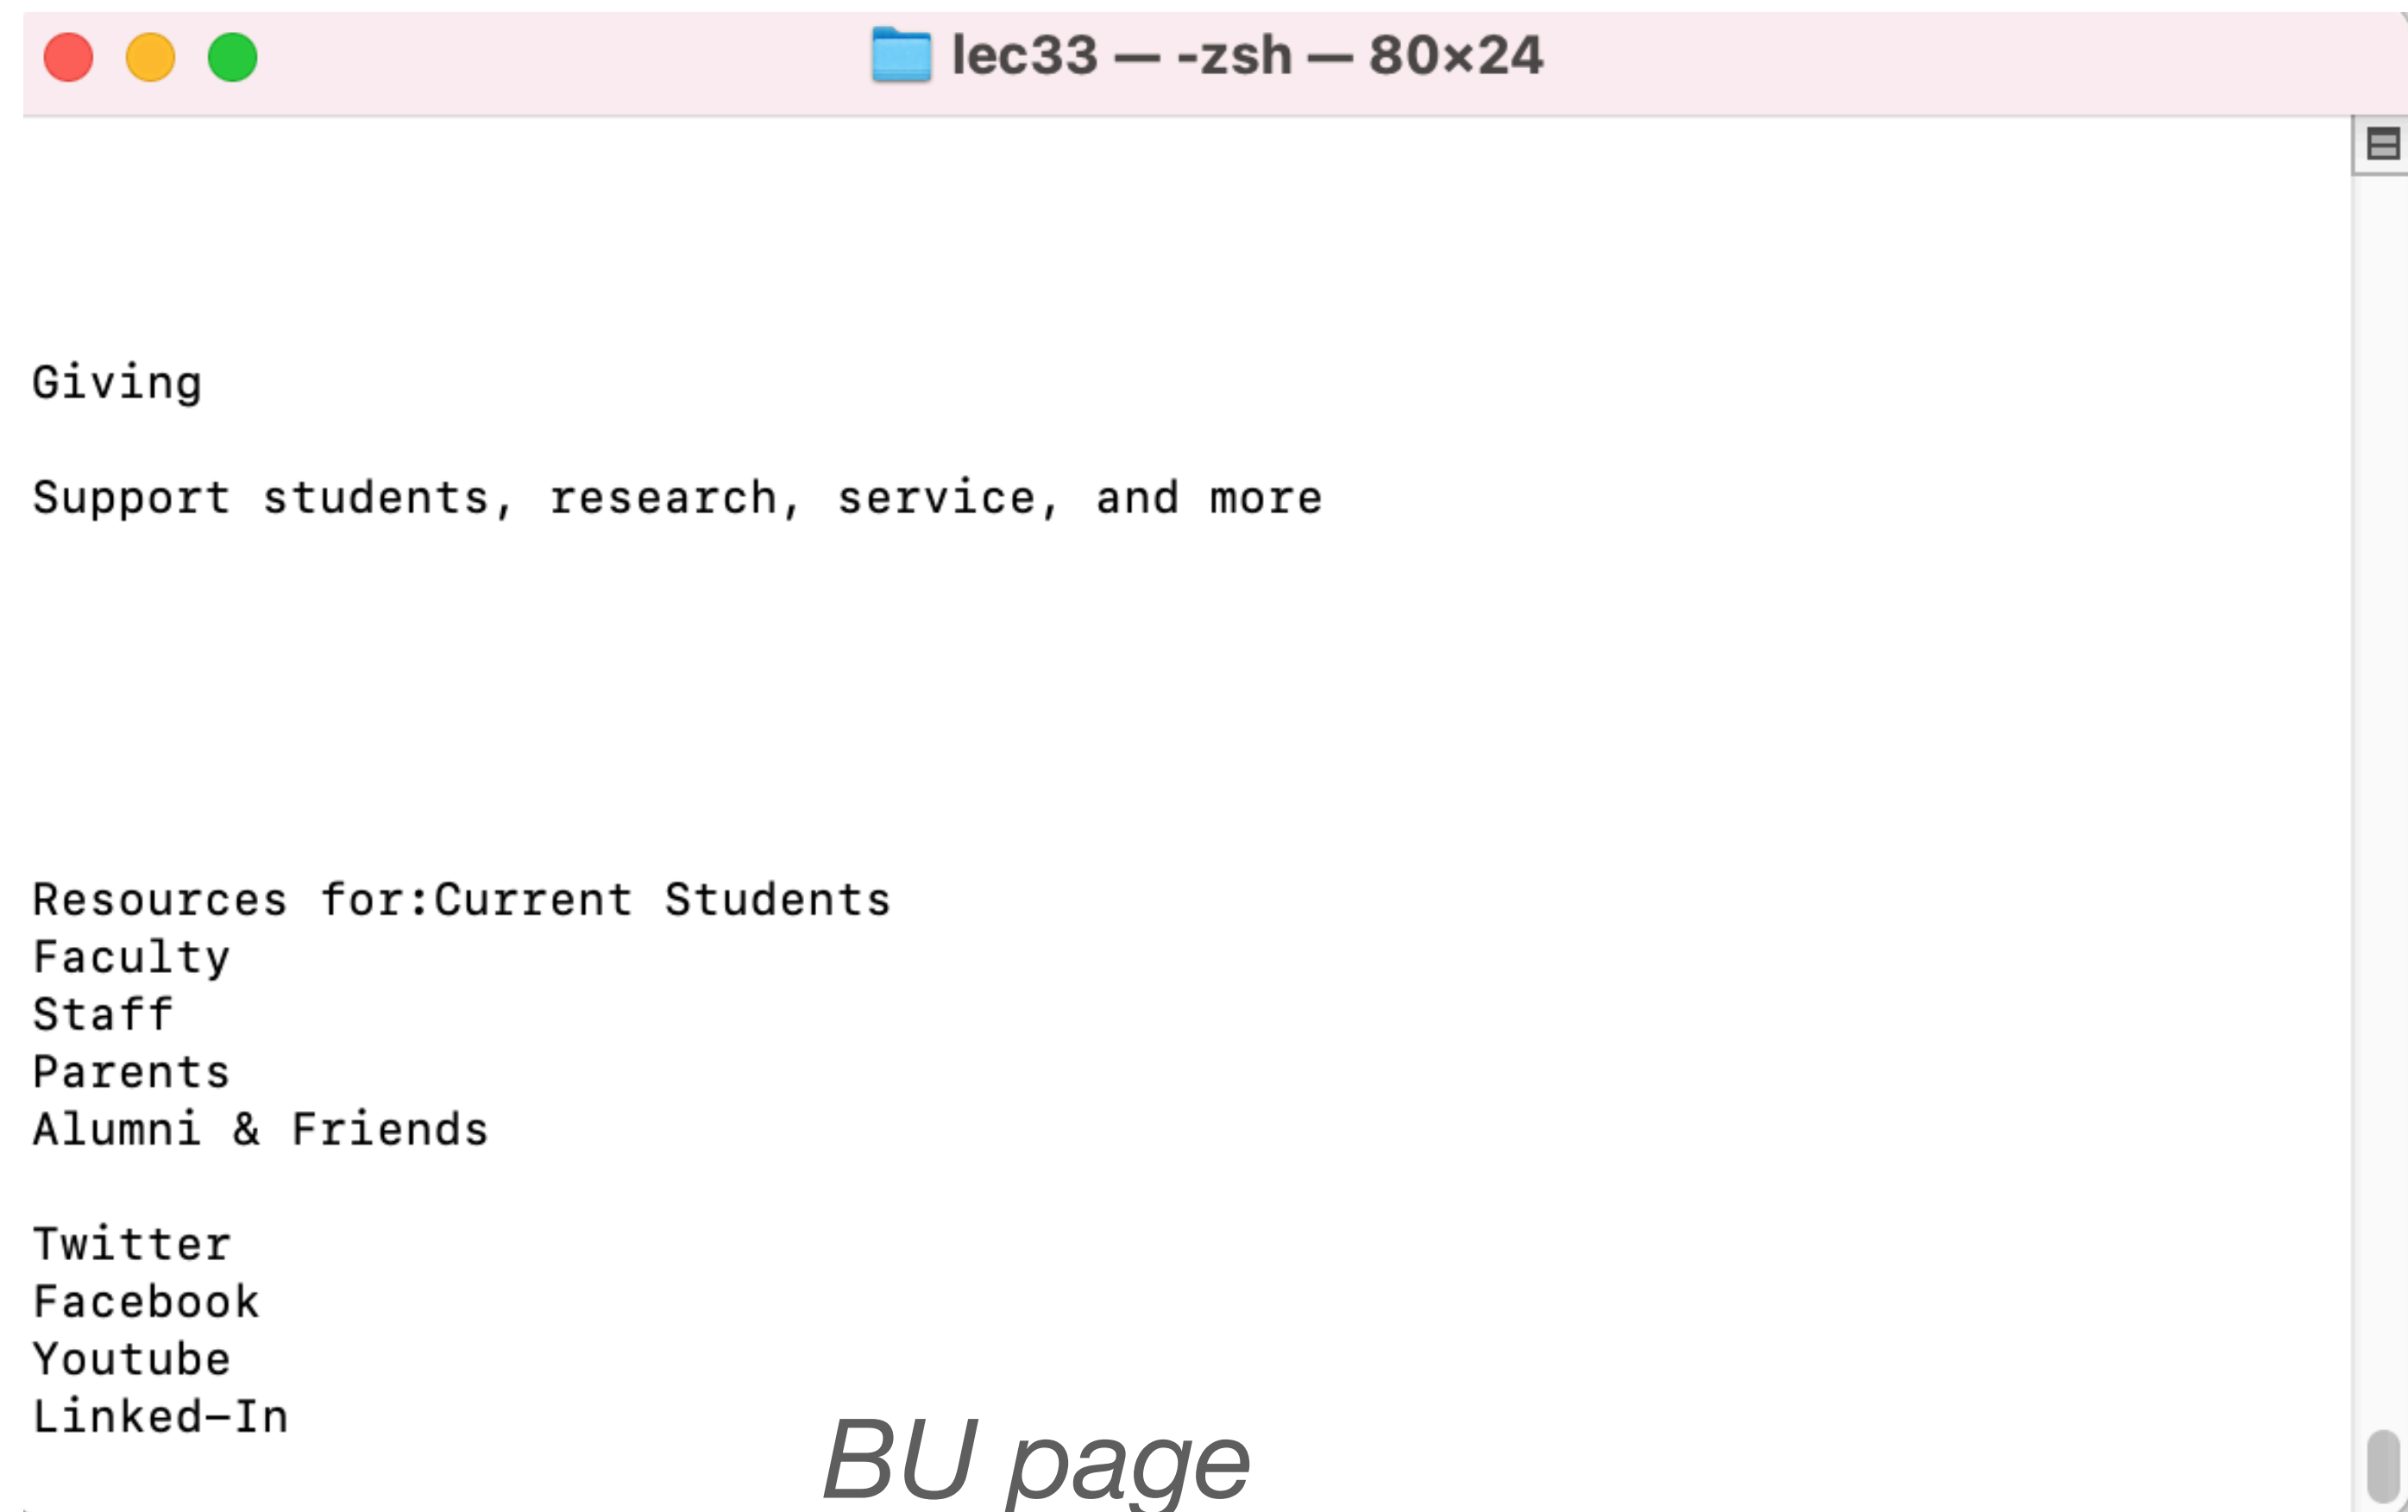
```
for link in soup.find_all('a'):  
    print(link.get('href'))
```

- Sample applications:
 - Find highly interconnected networks of pages that form "communities" to understand a pages ecosystem
 - Run PageRank, Google's algorithm that determines importance of page based on who links to you



Finding the readable text

```
print(soup.get_text())
```



Finding the readable text

```
print(soup.get_text())
```

- Sample applications:
 - Train machine learning to classify page topics based on their text (for use in retrieval later, for example)

Do you **drive** a **Toyota Corolla**,
but wonder whether a **Honda Fit** would be **better**?
We **test-drove** **2022** models of both cars...

Finding the readable text

```
print(soup.get_text())
```

- Sample applications:
 - Train machine learning to classify page topics based on their text
 - Determine with "sentiment analysis" whether web content is positive or negative

I **love** this Dyson vacuum

I **hate** Dysons so much

Finding the readable text

```
print(soup.get_text())
```

- Sample applications:
 - Train machine learning to classify page topics based on their text
 - Determine with "sentiment analysis" whether web content is positive or negative
 - Scan the text for facts to extract, like who starred in what film

Kevin Smith's latest film, Clerks 3, is a stunning achievement

Finding other things in the HTML besides links

- Find the page title: `soup.title.string`
- Find bigger headings on a page (<h1> through <h3>):

```
for heading in soup.find_all(['h1', 'h2', 'h3']):  
    print(heading.text)
```

- Search for strings containing a particular regular expression:

```
import re  
soup.find_all(string=re.compile("Dormouse"))
```

Finding div classes within a page

- Div tags are often used to divide up a page logically
- The different logical pieces of the page often have different **class** names that the CSS refers to in styling the page
- We can take advantage of this to extract just text from particular parts of the page we care about
- ```
divs = soup.find_all('div', class_='your_classname_here')
for div in divs:
 print(div.get_text())
```

# Example: Gamespot review titles

Latest Reviews

Sort By: Release Date

Search Reviews

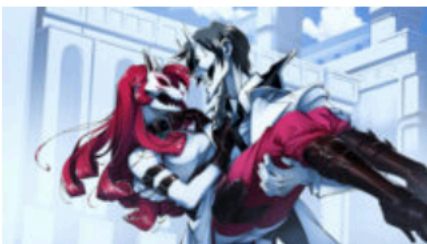
PC

Neon White Review – Heavenly Sprint

🕒 6 days ago    💬 2    👍 2

9

Superb



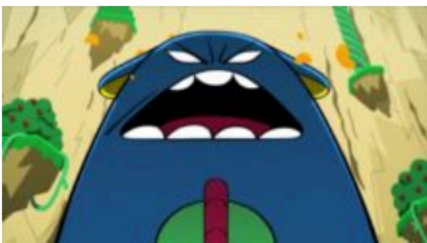
IOS

Poinpy Review - Moving On Up

🕒 12 days ago    💬 3    👍 4

8

Great



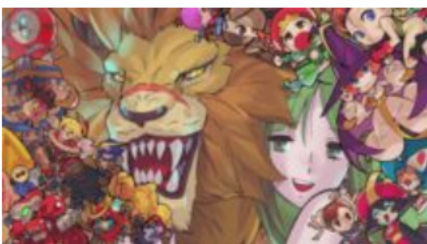
PS4

Capcom Fighting Collection Review - Family Feud

🕒 12 days ago    💬 1    👍 5


7

Good



Make your memories pop

Galaxy S22 Ultra with Google Photos



Find Reviews

Platform

Current platforms


Genre

All genres

Release Date

All-time

+ Advanced Options



NOW STREAMING

HBOMAX

SIGN UP NOW

PLANS START AT \$9.99/MONTH

- Site looks like the above
- We want the text in the review cards

# Example: Inspecting the page HTML

```
<div id="js-sort-filter-results">
```

```
 <section class="editorial river ">
```

```
 <div class='card-item base-flexbox flexbox-align-center width-100 border-bottom-grayscale--thin "><div class="card-item__img (
```

```
 <div class='card-item base-flexbox flexbox-align-center width-100 border-bottom-grayscale--thin "><div class="card-item__img (
```

- Inspection of the HTML and a little experimentation reveals that "card-item" is a class shared by each boxed review



# Example: Code to extract the review card text

```
import requests
from bs4 import BeautifulSoup
from textblob import TextBlob
import nltk
nltk.download('punkt')
import re

page = requests.get('https://www.gamespot.com/games/reviews/')
contents = page.content

soup = BeautifulSoup(contents)
reviews = soup.find_all('div', class_ = 'card-item')

for review in reviews:
 print(review.get_text())
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
PCNeon White Review – Heavenly Sprint7 days ago229Superb
IOSPoinpy Review – Moving On Up12 days ago348Great
PS4Capcom Fighting Collection Review – Family Feud12 days ago157Good
IOSDisney Mirrorverse Review – Shattered Dreams13 days ago1934Poor
PS4Sonic Origins Review – Going Fast, Again14 days ago1747Good
```

## Latest Reviews

Sort By: Release Date

Search Reviews

PC

### Neon White Review – Heavenly Sprint

6 days ago 2 2

9

Superb




IOS

### Poinpy Review - Moving On Up

12 days ago 3 4

8

Great




PS4

### Capcom Fighting Collection Review - Family Feud

12 days ago 1 5

7

Good



Make your  
memories pop

Galaxy S22 Ultra  
with Google Photos





# Example: Refining the search

- We could look for a div or other tag that more precisely captures the information we want

```
<h4 class="card-item__title">Neon White Review – Heavenly Sprint</h4>
```

```
soup = BeautifulSoup(contents)
reviews = soup.find_all('h4')
```

```
*Neon White Review – Heavenly Sprint
*Poinpy Review – Moving On Up
*Capcom Fighting Collection Review – Family Feud
*Disney Mirrorverse Review – Shattered Dreams
*Sonic Origins Review – Going Fast, Again
*Diablo Immortal Review – Evil On The Go
```

# Example: Using regular expressions

- We could alternately, or in addition, use regular expressions to clean and parse the information we got

```
page = requests.get('https://www.gamespot.com/games/reviews/')
contents = page.content

soup = BeautifulSoup(contents)
reviews = soup.find_all('h4')

for review in reviews:
 pattern = '(.)+ Review (.)+'
 result = re.search(pattern, review.get_text())
 if result:
 print(result.group(1))
```

```
Neon White
Poinpy
Capcom Fighting Collection
Disney Mirrorverse
Sonic Origins
Diablo Immortal
Fire Emblem Warriors: Three Hopes
Teenage Mutant Ninja Turtles: Shredder's Revenge
The Quarry
Roller Champions
Halo Infinite Multiplayer
Soundfall
Apex Legends Mobile
Hatsune Miku Project Diva Megamix+
Sniper Elite 5
```

# Summary

- We can pull text from either webpages to power machine learning
- We can be interested in the text, the graph, or both
- requests plus BeautifulSoup easily pulls text from websites
- It can take some trial and error to find the right patterns to search for