


Graphs and Centrality

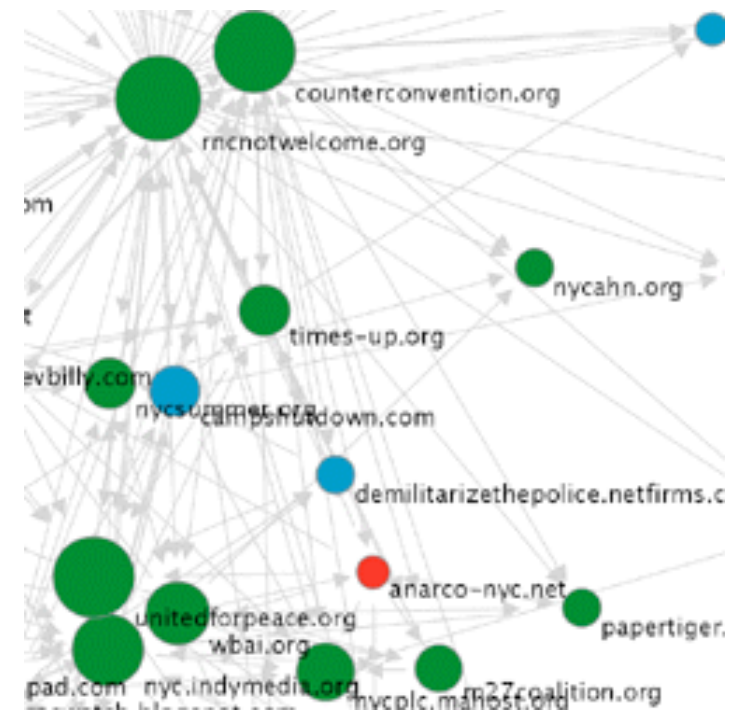
Professor Kevin Gold

The Graph Concept

- A graph consists of things (**vertices**) and connections between those things (**edges**)
- The connections can be asymmetric (directed graph) or symmetric (undirected graph)
 - Asymmetric/directed: Webpage A links to webpage B, but B doesn't link to A
 - Symmetric/undirected: We are Facebook friends — it's symmetric



Undirected graph with 4 vertices, 5 edges

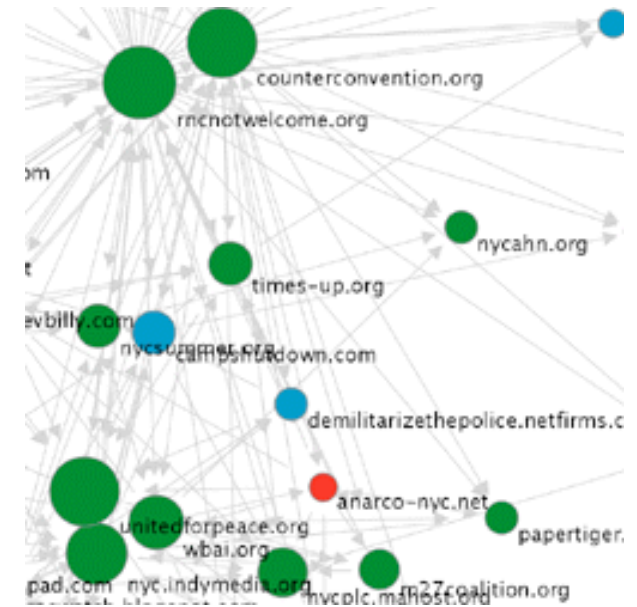


Piece of a directed graph of the web

web graph source: <http://farrall.org/webgraph/home.html>


Graph Algorithms Have Many Applications

- The World Wide Web - analyze structure to find important pages (PageRank)
- Social networks - Try to find influential individuals
- Computer networks - Find a path for data packets
- Pathfinding - Google Maps
- Graph databases (Neo4j) - for data structured like a graph
- Classic AI - treat puzzles as graphs to pathfind on
- And more!




More Specifically, a Graph is...

- A graph G consists of a set of vertices V (also called **nodes**) and a set of edges, E
 - The *number* of vertices and edges is referred to as $|V|$ and $|E|$
 - In an *undirected* graph, the edges are two-element subsets of V : $\{u, v\}$. No arrows on the drawing.
 - In a *directed* graph, the edges are ordered pairs of vertices (u, v) . We draw an arrow from u to v .
- The **degree** of a vertex is the number of edges it touches ("**in-degree**" and "**out-degree**" for directed)




More Specifically, a Graph is...

- Generally assume no self-loops
- Usually not multiple edges connecting the same two vertices — (u,v) is in the graph or not
- This would be called a “multigraph” (unusual)



Paths and Cycles

- “Is there a way to get from this vertex to that one?”
- A *path* from vertex A to vertex B is a sequence of vertices that leads from A to B — edges connect each vertex (going the right direction, if directed)
 - If all vertices distinct, it is a “simple path”
- A *cycle* is a path from a vertex to itself that contains at least one other vertex and doesn't repeat other vertices




A,B,D,C: simple path from A to C

A,B,D,A,B,D,C: non-simple path from A to C


A,B,D,A: cycle

Connected Graphs


- An undirected graph is **connected** if, for every pair of nodes u and v , there is a path from u to v .
- In directed graphs, this has two varieties.
 - A directed graph is **weakly connected** if it would be connected if it were undirected (the paths can “go the wrong way”)
 - A directed graph is **strongly connected** if there is a path (obeying the directions) from every node u to every other node v .
- Some algorithms assume graphs are connected in one of these ways.



not connected (B,C)




weakly connected (C)




strongly connected

Weights



unweighted graph




weighted graph


- By default, graphs are “unweighted”: the distance between two **neighbors** (nodes connected by an edge) is assumed to be 1.
 - And the *shortest path* between two nodes is the path between those nodes that uses the *fewest edges*
- **Weighted graphs** can represent distances and costs of paths between vertices (each edge gets a number that is its weight)
 - Shortest path is path with smallest sum of weights.

Trees

- A **tree** is any connected graph that has no cycles. It doesn't necessarily need to be hierarchical.
- You can choose one node to be the **root** — grab the tree and let the rest of the graph “hang” from that.
 - Each node besides the root will have one neighbor closer to the root — its **parent**
 - ...So a tree has exactly how many edges?
 - The **depth** of a node is its distance from the root



Totally a tree




Same tree, if A is root

Who's Important in this Network? **Centrality**

- Data science task: here is a graph of a social network; who is most influential?
- Influence could come in several forms:
 - Degree centrality - most direct friends
 - Closeness centrality - just a few hops from anyone
 - Betweenness centrality - a "gatekeeper" between subgraphs
 - Eigenvector centrality - considered important by important nodes

Degree Centrality


- Simply the degree (number of neighbors) of the vertex
- If graph is directed, could focus on in-degree or out-degree as appropriate
- Example: someone with a lot of Twitter followers is important
 - A little easy to game with bots and fake accounts



High in-degree centrality

Degree Centrality

Twitter accounts with the most followers worldwide as of April 2022 (in millions)



Source
Socialtracker
© Statista 2022


Additional Information:
Worldwide; Socialtracker; April 2022

Closeness Centrality

- The inverse of the average number of steps in a closest path to each other node

- $C(x) = \frac{N - 1}{\sum_y d(y, x)}$ where $N-1$ is a count of all the other nodes besides x and $d(y, x)$ is a shortest path distance

- Larger for nodes that are just a few hops away from the whole graph - a broadcast message has "reach"




Closeness $11/(7*1+4*2)=0.73$

Betweenness Centrality

- Finds nodes that connect one community to another
- Betweenness of v is sum over start and end locations of proportion of shortest paths that pass through v :

$$\sum_{s \neq v \neq t} \frac{\sigma_{svt}}{\sigma_{st}} \text{ where } \sigma_{svt} \text{ is a count of shortest paths}$$


from s to t through v and σ_{st} is a count of shortest paths from s to t (maybe including v , maybe not)



*s, t on opposite sides of v
contribute 1, on same side
contribute 0*

Eigenvector Centrality

- A term for what Google's PageRank does
- With some linear algebra, calculates the proportion of time a random walk through the graph would pass through any particular node
- Being linked to by popular sites does more to increase your popularity than links from small sites - thus harder to "game"





Finding Shortest Paths: Breadth-First Search

- Closeness and betweenness centrality both need to find shortest paths to calculate them
- There are two fundamental algorithms for finding a path from one node to another on an **unweighted** graph: **breadth-first search** and **depth-first search**
- But breadth-first search finds the *shortest* path, and depth-first search generally doesn't.
- General BFS idea: fan out from the start, methodically trying all closer vertices before all farther ones

Breadth-First Search


- Basic idea: Visit all nodes one step away, then all nodes two steps away, then ...
- In exploring, we can build a tree, where the nodes at depth i are i hops away from the start using a shortest path
- We can also use a **queue** to track which nodes to explore next — a **first-in, first out** list where new elements “go to the end of the line”



Breadth-First Search

- Add start node **s** to the **queue** of nodes to explore
- Initialize “**discovered**” set, tracking which nodes have been seen before, to **s**
- Initialize distance table with $d[s] = 0$ and create an empty parent table (representing the tree)
- While the queue of nodes to explore is not empty:
 - Pull off the node p at the **front** of the queue (oldest node in queue)
 - For each neighbor n of p :
 - If n is not in discovered set:
 - Add p as n 's parent in parent table; set $d[n]$ to $d[p]+1$
 - If n is the target **t**, return the path to **t** (following the parent table from **t** back to **s**)
 - Add n to the **end** of the queue and add to “discovered” set
- Return "no path"


BFS Example




Queue
A

Tree
A


BFS Example



Queue
B, H




BFS Example




Queue
H,C


Tree




BFS Example




Queue
C




BFS Example




Queue
D, I




BFS Example




Queue
I, E




BFS Example



Queue
E




BFS Example




Queue
F,J

Tree




BFS Example




Queue
J


Tree



BFS Example




Queue
~empty~



*Follow the
parent links
from G to
get the path*

More on Queues


- In an efficiently implemented queue, adding new elements to the back or pulling elements from the front is fast.
- This is possible with 2 additional pointers (or references) in a linked list, “First” and “Last.” Last, in particular, allows direct access to the end.
- In Python, a **deque** data structure offers the quick access to the front and back that a basic list lacks
 - In basic list, only append is fast; dequeue requires shifting elements



Graph Representations

- Two approaches:
 - *Adjacency matrix* : A $|V| \times |V|$ array that is 1 at u,v if there is an edge from u to v
 - *Adjacency list*: $|V|$ lists of vertices, where each list contains the vertices a particular vertex has an edge to
 - In Python, can be a dictionary of lists

Examples of graph representations



Adjacency Matrix

	A	B	C	D
A	0	1	0	1
B	0	0	1	0
C	0	0	0	0
D	0	0	1	0

Adjacency Lists


A: [B, D]
B: [C]
C: []
D: [C]

Graph Representation

Tradeoffs: Adjacency Matrix

- Pro: Direct lookup of connectivity of two vertices
- Cons:
 - V^2 space when there aren't necessarily that many edges
 - Need to iterate through V entries to check all neighbors of a node

	A	B	C	D
A	0	1	0	1
B	0	0	1	0
C	0	0	0	0
D	0	0	1	0



Graph Representation

Tradeoffs: Adjacency Lists

- Pros:
 - Checking all edges will not waste time on non-existent edges
 - Takes less memory
- Con: Checking connectivity of a particular pair of vertices requires traversing a list


Adjacency Lists

A: [B, D]

B: [C]


C: []

D: [C]



BFS With Adjacency Matrices is Slower


- With adjacency matrices, in adding all neighbors of a node to queue, need to check all entries of the matrix for that node's row
- That could be many more operations than just checking the edges that exist



	A	B	C	D
A		1		1
B				
C				
D			1	

Calculating Centrality With BFS

- Closeness centrality requires 1 BFS with no "target" to find shortest path lengths to all nodes
- Betweenness centrality can be estimated by choosing random s, t pairs and running BFS
- Degree and eigenvector centrality don't need BFS



Summary — Graphs

- Graphs are an extremely versatile way of abstractly reasoning about relationships and connections
- One thing we can do with a social network graph is calculate different kinds of centrality - find the most important people
- Breadth-first search is an efficient tool for finding shortest paths on graphs - which can be used for centrality